

Session 1: HPC and the Julia framework

OBJECTIVE: Confirm Julia framework and Base speed

KR1: Use `@code_*` to examine a simple function. The `*` is replaceable by `native`, `typed`, `warntype`, and others. Discover them.

First, I want to look at the docs of each macro function:

```
In [1]: ?code_native
```

```
search: code_native @code_native
```

Out[1]: `code_native([io=stdout,], f, types; syntax=:att, debuginfo=:default)`
Prints the native assembly instructions generated for running the method matching the given generic function and type signature to `io`. Switch assembly syntax using `syntax` symbol parameter set to `:att` for AT&T syntax or `:intel` for Intel syntax. Keyword argument `debuginfo` may be one of `source` (default) or `none`, to specify the verbosity of code comments.

```
In [2]: @code_native 10 / 2
```

```
        .section          __TEXT,__text,regular,pure_instructions
; | @ int.jl:93 within `/'
; | | @ float.jl:206 within `float'
; | | | @ float.jl:191 within `AbstractFloat'
; | | | | @ float.jl:94 within `Float64'
; | | | |    vcvtsi2sd      %rdi, %xmm0, %xmm0
; | | | |    vcvtsi2sd      %rsi, %xmm1, %xmm1
; | | LLL
; | | @ int.jl:93 within `/' @ float.jl:335
; | |    vdivsd    %xmm1, %xmm0, %xmm0
; | | @ int.jl:93 within `/'
; | |    retq
; | @ int.jl:93 within `<invalid>'
; |    nop
; L
```

```
In [3]: ?code_typed
```

```
search: code_typed @code_typed code_warntype @code_warntype
```

Out[3]: `code_typed(f, types; optimize=true, debuginfo=:default)`
Returns an array of type-inferred lowered form (IR) for the methods matching the given generic function and type signature. The keyword argument `optimize` controls whether additional optimizations, such as inlining, are also applied. The keyword `debuginfo` controls the amount of code metadata present in the output, possible options are `:source` or `:none`.

```
In [4]: @code_typed 10 / 2
```

```
Out[4]: CodeInfo(
1 - %1 = Base.sitofp(Float64, x)::Float64
  |   %2 = Base.sitofp(Float64, y)::Float64
  |   %3 = Base.div_float(%1, %2)::Float64
  └─ return %3
) => Float64
```

```
In [5]: ?code_warntype
```

```
search: code_warntype @code_warntype
```

```
Out[5]: code_warntype([io::IO], f, types; debuginfo=:default)
```

Prints lowered and type-inferred ASTs for the methods matching the given generic function and type signature to `io` which defaults to `stdout`. The ASTs are annotated in such a way as to cause "non-leaf" types to be emphasized (if color is available, displayed in red). This serves as a warning of potential type instability. Not all non-leaf types are particularly problematic for performance, so the results need to be used judiciously. In particular, unions containing either `missing` or `nothing` are displayed in yellow, since these are often intentional.

Keyword argument `debuginfo` may be one of `:source` or `:none` (default), to specify the verbosity of code comments.

See `@code_warntype` for more information.

```
In [6]: @code_warntype 10 / 2
```

```
Variables
```

```
#self#::Core.Const{}
```

```
x::Int64
```

```
y::Int64
```

```
Body::Float64
```

```
1 - %1 = Base.float(x)::Float64
```

```
| %2 = Base.float(y)::Float64
```

```
| %3 = (%1 / %2)::Float64
```

```
|_ return %3
```

```
In [7]: ?code_llvm
```

```
search: code_llvm @code_llvm
```

```
Out[7]: code_llvm([io=stdout,], f, types; raw=false, dump_module=false, optimize=true, debuginfo=:default)
```

Prints the LLVM bitcodes generated for running the method matching the given generic function and type signature to `io`.

If the `optimize` keyword is unset, the code will be shown before LLVM optimizations. All metadata and `dbg.*` calls are removed from the printed bitcode. For the full IR, set the `raw` keyword to true. To dump the entire module that encapsulates the function (with declarations), set the `dump_module` keyword to true. Keyword argument `debuginfo` may be one of `source` (default) or `none`, to specify the verbosity of code comments.

```
In [8]: @code_llvm 10 / 2
```

```
; @ int.jl:93 within `/'
define double @"julia/_1801"(i64 signext %0, i64 signext %1) {
top:
; @ float.jl:206 within `float'
; | @ float.jl:191 within `AbstractFloat'
; || @ float.jl:94 within `Float64'
    %2 = sitofp i64 %0 to double
    %3 = sitofp i64 %1 to double
```

```

; LLL
; @ int.jl:93 within `/' @ float.jl:335
%4 = fdiv double %2, %3
; @ int.jl:93 within `/'
ret double %4
}

```

In [9]: `@code_llvm 10.0 / 2.0`

```

; @ float.jl:335 within `/'
define double @"julia_/_1814"(double %0, double %1) {
top:
    %2 = fdiv double %0, %1
    ret double %2
}

```

What's interesting here is under the hood, if the inputs for division are not floats/doubles they're converted first.

Personally, among these macros `@code_warntype` is likely the one I'm going to use the most in the future.

KR2: Demonstrate that Julia is able to determine constants in codes.

Here we can see that the euler constant is a constant inside the function (0x4005BF0A8B145769).

In [10]: `f(x) = log(x) * e`

Out[10]: `f (generic function with 1 method)`

In [11]: `@code_llvm f(1.0)`

```

; @ In[10]:1 within `f'
define double @julia_f_1870(double %0) {
top:
    %1 = call double @j_log_1872(double %0)
;  └ @ promotion.jl:322 within `*' @ float.jl:332
    %2 = fmul double %1, 0x4005BF0A8B145769
;  └
    ret double %2
}

```

If we don't have a constant, we can see that it won't reflect in `@code_llvm`

In [12]: `g(x) = log(x) * x`

Out[12]: `g (generic function with 1 method)`

In [13]: `@code_llvm g(1.0)`

```

; @ In[12]:1 within `g'
define double @julia_g_1897(double %0) {
top:
    %1 = call double @j_log_1899(double %0)
;  └ @ float.jl:332 within `*'
    %2 = fmul double %1, %0
;  └
    ret double %2
}

```

KR3: Demonstrate Julia's type-inference and multiple dispatch.

Following demo in class, to demonstrate multiple dispatch

```
In [14]: what(x) = "$ (x) is a $(typeof(x))" #most general case
          what(x::String) = "$ (x) is a string"
          what(x::Number) = "$ (x) is a number"
          what(x::Real) = "$ (x) is real number"
          what(x::Rational) = "$ (x) is rational number"
          what(x::Int) = "$ (x) is an integer"
          what(x::Complex) = "$ (x) is complex number"
```

Out[14]: what (generic function with 7 methods)

```
In [15]: ?what
```

search: **what** Cwchar_t

Out[15]: No documentation found.

what is a Function .

```
# 7 methods for generic function "what":
[1] what(x::String) in Main at In[14]:2
[2] what(x::Rational) in Main at In[14]:5
[3] what(x::Int64) in Main at In[14]:6
[4] what(x::Real) in Main at In[14]:4
[5] what(x::Complex) in Main at In[14]:7
[6] what(x::Number) in Main at In[14]:3
[7] what(x) in Main at In[14]:1
```

Just to try a lot of things:

```
In [16]: for x in [π, 1, 1.0, "x", 'x', 1+0im,
                  [1,2], ["x", "y"], ['x', 'y'], ["x", 'y'], [Int64(1), 1.0]]
          println(what(x))
        end
```

```
π is real number
1 is an integer
1.0 is real number
x is a string
x is a Char
1 + 0im is complex number
[1, 2] is a Vector{Int64}
["x", "y"] is a Vector{String}
['x', 'y'] is a Vector{Char}
Any["x", 'y'] is a Vector{Any}
[1.0, 1.0] is a Vector{Float64}
```

Vectors having specific typings that can be specified is interesting having never thought about this when I was primarily using Python before.

One other thing I have to wrap my head around in Julia is the subset of types. For example π in the previous output is a real number (which it is), but when we specifically call `typeof` on it, it's more specific(?):

Irrational .

```
In [17]: typeof(π)
```

Out[17]: Irrational{::π}

KR4: Show the difference, if any, between your own sum function `my_sum(x::Vector)` and `@time` . Use a `for`-loop for your *customized* sum function.

Unrelated to the KR but I realized here that Σ \Sigma is a different character here from Σ \Sum

```
In [18]: """
           my_sum( x )
Return the sum of all elements in an input vector `x`
- Input: `x::Vector`
- Output: `Σ::Number`
           """
           function my_sum(x::Vector)
               Σ = zero(eltype(x))
               for i in x
                   Σ += i
               end
               return Σ
           end
```

Out[18]: `my_sum`

```
In [19]: ?my_sum
```

search: `my_sum`

Out[19]: `my_sum(x)`
Return the sum of all elements in an input vector `x`

- Input: `x::Vector`
- Output: `Σ::Number`

```
In [20]: ?sum
```

search: `sum` `sum!` `summary` `cumsum` `my_sum` `cumsum!` `isnumeric` `VersionNumber`

Out[20]: `sum(f, itr; [init])`
Sum the results of calling function `f` on each element of `itr` .

The return type is `Int` for signed integers of less than system word size, and `UInt` for unsigned integers of less than system word size. For all other arguments, a common return type is found to which all arguments are promoted.

The value returned for empty `itr` can be specified by `init` . It must be the additive identity (i.e. zero) as it is unspecified whether `init` is used for non-empty collections.

!!! compat "Julia 1.6" Keyword argument `init` requires Julia 1.6 or later.

Examples

```
jldoctest
julia> sum(abs2, [2; 3; 4])
29
```

Note the important difference between `sum(A)` and `reduce(+, A)` for arrays with small integer eltype:

```
jldoctest
julia> sum(Int8[100, 28])
128
```

```
julia> reduce(+, Int8[100, 28])
-128
```

In the former case, the integers are widened to system word size and therefore the result is 128. In the latter case, no such widening happens and integer overflow results in -128.

```
sum(itr; [init])
```

Returns the sum of all elements in a collection.

The return type is `Int` for signed integers of less than system word size, and `UInt` for unsigned integers of less than system word size. For all other arguments, a common return type is found to which all arguments are promoted.

The value returned for empty `itr` can be specified by `init`. It must be the additive identity (i.e. zero) as it is unspecified whether `init` is used for non-empty collections.

!!! compat "Julia 1.6" Keyword argument `init` requires Julia 1.6 or later.

Examples

```
jldoctest
julia> sum(1:20)
210
```

```
julia> sum(1:20; init = 0.0)
210.0
```

```
sum(A::AbstractArray; dims)
```

Sum elements of an array over the given dimensions.

Examples

```
jldoctest
julia> A = [1 2; 3 4]
2×2 Matrix{Int64}:
 1  2
 3  4
```

```
julia> sum(A, dims=1)
1×2 Matrix{Int64}:
 4  6
```

```
julia> sum(A, dims=2)
2×1 Matrix{Int64}:
 3
 7
```

```
sum(f, A::AbstractArray; dims)
```

Sum the results of calling function `f` on each element of an array over the given dimensions.

Examples

```
jldoctest
julia> A = [1 2; 3 4]
2×2 Matrix{Int64}:
 1  2
 3  4

julia> sum(abs2, A, dims=1)
1×2 Matrix{Int64}:
10 20

julia> sum(abs2, A, dims=2)
2×1 Matrix{Int64}:
 5
25
```

Custom summation function is twice as slower as the base function

```
In [21]: @time my_sum(rand(1000000))

0.009748 seconds (13 allocations: 7.649 MiB)
Out[21]: 499928.0970995862
```

```
In [22]: @time sum(rand(1000000))

0.002831 seconds (2 allocations: 7.629 MiB)
Out[22]: 499960.4828240783
```

KR5: Replicate plotting the Mandelbrot. Use a separate file `Mandelbrot.jl` to contain the function code. Use `include()` function to load the file.

Just basing the code off of our textbook

```
In [23]: ;cat Mandelbrot.jl

"""
    mandel(c)
Given a complex number, computes whether after a certain number of iterations
`f_c(z) = z^2 + c` converges or not.
"""
function mandel(c)
    z = c
    maxiter = 80
    for n in 1:maxiter
        if abs(z) > 2
            return n - 1
        end
        z = z^2 + c
    end
    return maxiter
end

"""
    mandel_grid( xrange::Tuple{Float64,Float64}, yrange::Tuple{Float64,Float64}; n=100 )
```

Applies the `mandel()` function over a grid with set xrange, yrange with shape `n` x `n``.

Arguments:

xrange (Tuple{Float64, Float64}): bounds of the grid along x. Defaults to (-1.0, 1.0)
yrange (Tuple{Float64, Float64}): bounds of the grid along y. Defaults to (-1.0, 1.0)
n (Int64): Dimensions of the grid (length of one side). Defaults to 100.

"""

```
function mandel_grid(  
    xrange::Tuple{Float64,Float64}=(-1.0, 1.0),  
    yrange::Tuple{Float64,Float64}=(-1.0, 1.0),  
    n::Int64=100,  
)  
    grid = zeros(n, n)  
    xval = range(xrange[1], xrange[2]; length=n)  
    yval = range(yrange[1], yrange[2]; length=n)  
    for i in 1:n, j in 1:n  
        grid[i, j] = mandel(xval[i] + im * yval[j])  
    end  
    return grid  
end
```

In [24]: `include("Mandelbrot.jl")`

Out[24]: `mandel_grid`

In [25]: `?mandel`

search: `mandel mandel_grid`

Out[25]: `mandel(c)`

Given a complex number, computes whether after a certain number of iterations $f_c(z) = z^2 + c$ converges or not.

In [26]: `?mandel_grid`

search: `mandel_grid`

Out[26]: `mandel_grid(xrange::Tuple{Float64,Float64}, yrange::Tuple{Float64,Float64}; n=100)`

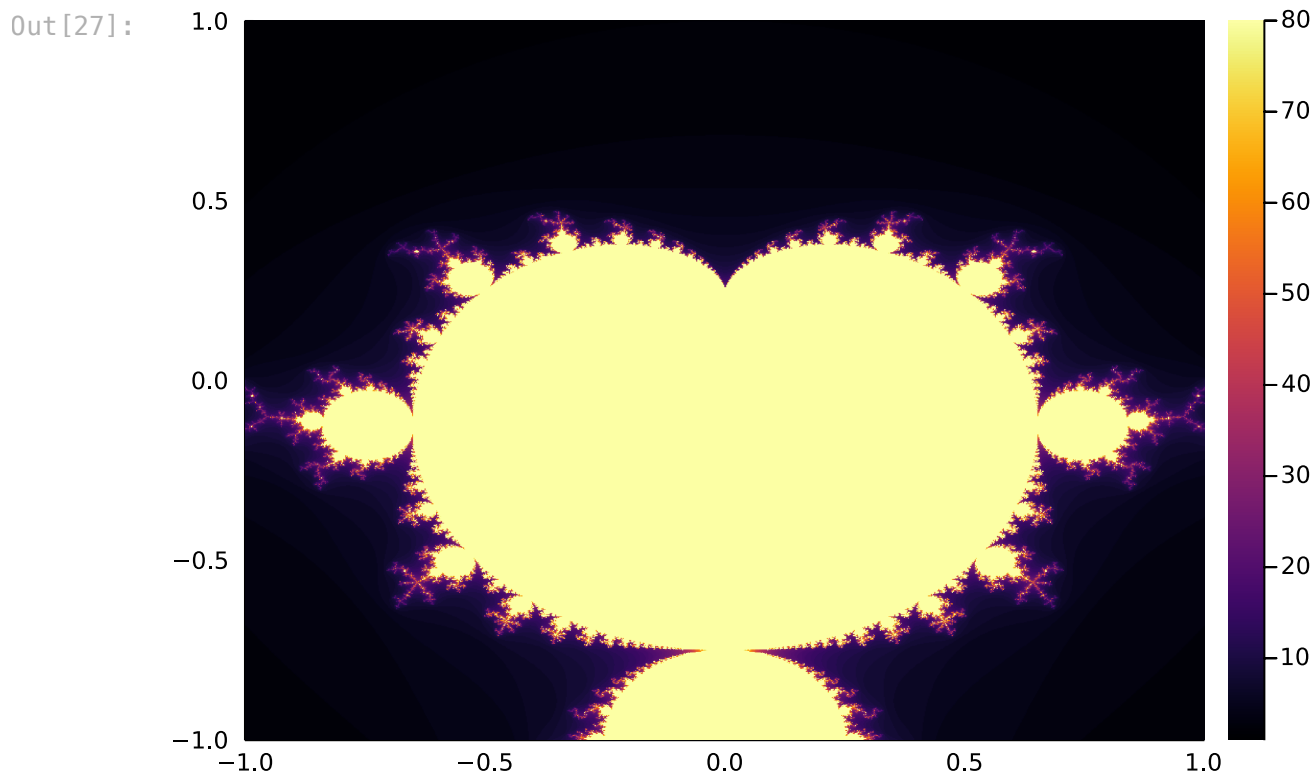
Applies the `mandel()` function over a grid with set xrange, yrange with shape `n x n`.

Arguments: xrange (Tuple{Float64, Float64}): bounds of the grid along x. Defaults to (-1.0, 1.0) yrange (Tuple{Float64, Float64}): bounds of the grid along y. Defaults to (-1.0, 1.0) n (Int64): Dimensions of the grid (length of one side). Defaults to 100.

I chose to create a function to generate the arrays in the grid and return them separate from just plotting them so that when we profile / benchmark the functions plotting won't reflect in the runtime.

In [27]: `using Plots`

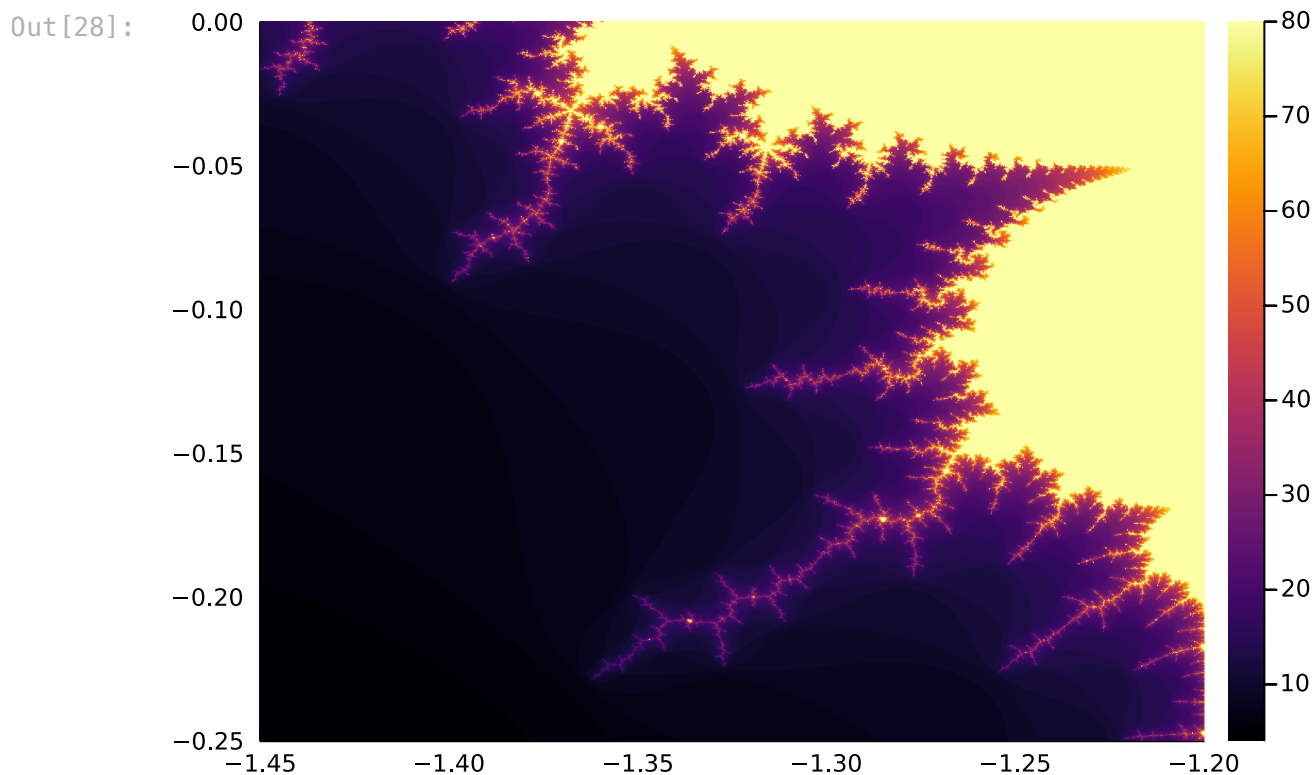
```
n = 1000  
xrange = (-1.0, 1.0)  
yrange = (-1.0, 1.0)  
x = LinRange(xrange[1], yrange[2], n)  
y = LinRange(yrange[1], yrange[2], n)  
z = mandel_grid(xrange=xrange, yrange=yrange, n=n)  
  
heatmap(x, y, z)
```

Zooming in to a specific spot:

```
In [28]: n = 1000
xrange = (-1.45, -1.20)
yrange = (-0.25, 0.0)
x = LinRange(xrange[1], xrange[2], n)
y = LinRange(yrange[1], yrange[2], n)
z = mandel_grid(xrange=xrange, yrange=yrange, n=n)

heatmap(x, y, z)
```



KR6: Plot of the time it takes for the function to run using `@time` macro for the given grid size `n`.

I used `@elapsed` so I can extract the time in seconds since I can assign it to a variable.

```
In [29]: @elapsed mandel_grid()
```

```
Out[29]: 0.004806154
```

```
In [30]: @time mandel_grid()
```

```
Out[30]: 0.007881 seconds (2 allocations: 78.203 KiB)
100×100 Matrix{Float64}:
 2.0  2.0  2.0  2.0  2.0  2.0  2.0  2.0  ...  2.0  2.0  2.0  2.0  2.0  2.0  2.0
 2.0  2.0  2.0  2.0  2.0  2.0  2.0  2.0      2.0  2.0  2.0  2.0  2.0  2.0  2.0
 2.0  2.0  2.0  2.0  2.0  2.0  2.0  2.0      2.0  2.0  2.0  2.0  2.0  2.0  2.0
 2.0  2.0  2.0  2.0  2.0  2.0  2.0  2.0      2.0  2.0  2.0  2.0  2.0  2.0  2.0
 2.0  2.0  2.0  2.0  2.0  2.0  2.0  2.0      2.0  2.0  2.0  2.0  2.0  2.0  2.0
 2.0  2.0  2.0  2.0  2.0  2.0  2.0  2.0  ...  2.0  2.0  2.0  2.0  2.0  2.0  2.0
 2.0  2.0  2.0  2.0  2.0  2.0  2.0  2.0      2.0  2.0  2.0  2.0  2.0  2.0  2.0
 2.0  2.0  2.0  2.0  2.0  2.0  2.0  3.0      2.0  2.0  2.0  2.0  2.0  2.0  2.0
 2.0  2.0  2.0  2.0  2.0  2.0  3.0  3.0      3.0  2.0  2.0  2.0  2.0  2.0  2.0
 2.0  2.0  2.0  2.0  2.0  2.0  3.0  3.0      3.0  2.0  2.0  2.0  2.0  2.0  2.0
 2.0  2.0  2.0  2.0  2.0  3.0  3.0  3.0  ...  3.0  3.0  2.0  2.0  2.0  2.0  2.0
 2.0  2.0  2.0  2.0  3.0  3.0  3.0  3.0      3.0  3.0  3.0  2.0  2.0  2.0  2.0
 2.0  2.0  2.0  3.0  3.0  3.0  3.0  3.0      3.0  3.0  3.0  3.0  2.0  2.0  2.0
 ⋮      ⋮      ⋮      ⋮      ⋮      ⋮      ⋮      ⋮      ⋮      ⋮      ⋮      ⋮      ⋮      ⋮
 1.0  1.0  1.0  1.0  1.0  1.0  1.0  1.0      1.0  1.0  1.0  1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0  1.0  1.0  1.0  1.0      1.0  1.0  1.0  1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0  1.0  1.0  1.0  1.0  ...  1.0  1.0  1.0  1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0  1.0  1.0  1.0  1.0      1.0  1.0  1.0  1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0  1.0  1.0  1.0  1.0      1.0  1.0  1.0  1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0  1.0  1.0  1.0  1.0      1.0  1.0  1.0  1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0  1.0  1.0  1.0  1.0      1.0  1.0  1.0  1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0  1.0  1.0  1.0  1.0      1.0  1.0  1.0  1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0  1.0  1.0  1.0  1.0  ...  1.0  1.0  1.0  1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0  1.0  1.0  1.0  1.0      1.0  1.0  1.0  1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0  1.0  1.0  1.0  1.0      1.0  1.0  1.0  1.0  1.0  1.0  1.0
 1.0  1.0  1.0  1.0  1.0  1.0  1.0  1.0      1.0  1.0  1.0  1.0  1.0  1.0  1.0
```

I run gridsizes for differen 2^n values just to cover ground quicker.

```
In [31]: iter = 14
n_list = 2 .^ round.(Int64,
                    (LinRange(1, iter, iter)))
          )
```

```
Out[31]: 14-element Vector{Int64}:
 2
 4
 8
16
32
64
128
256
512
1024
2048
4096
8192
16384
```

```
In [32]:
```

```

elapsed = Array{Float64}(undef, iter)
for i in 1:iter
    elapsed[i] = @elapsed mandel_grid(n=n_list[i])
end
elapsed

```

```

Out[32]: 14-element Vector{Float64}:
 0.021363124
 1.4215e-5
 2.569e-5
 0.000111479
 0.000452867
 0.001795561
 0.008040005
 0.03193065
 0.114736369
 0.47119516
 2.062564557
 7.564541668
33.282219048
156.832726844

```

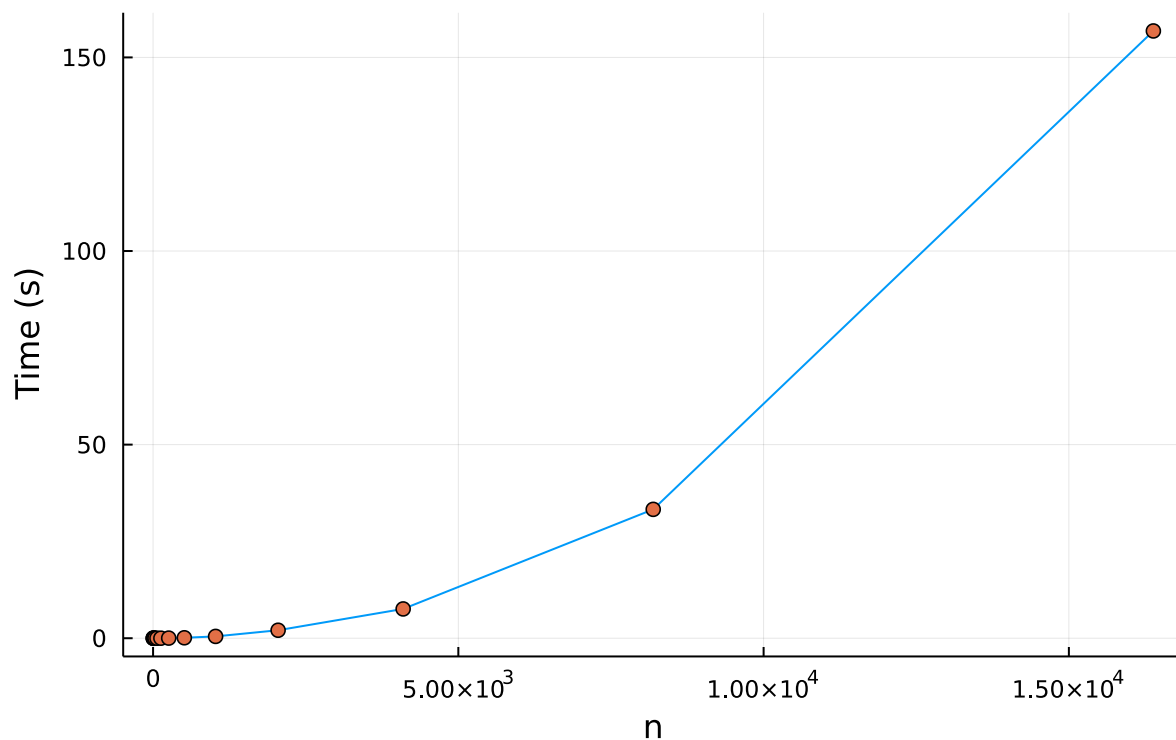
Plotting n vs time, it seems like the computational complexity is polynomial in time.

```

In [33]: plot(n_list, elapsed, label=false)
scatter!(n_list,
         elapsed,
         xlabel="n",
         ylabel="Time (s)",
         title="Running time for `mandel_grid()` for different gridsizes",
         label=false)

```

Out[33]: Running time for `mandel_grid()` for different gridsizes



However, I realized that the x-axis might probably not need to be n but n^2 . Squaring the x-axis, we see that it looks like it's actually linear.

```

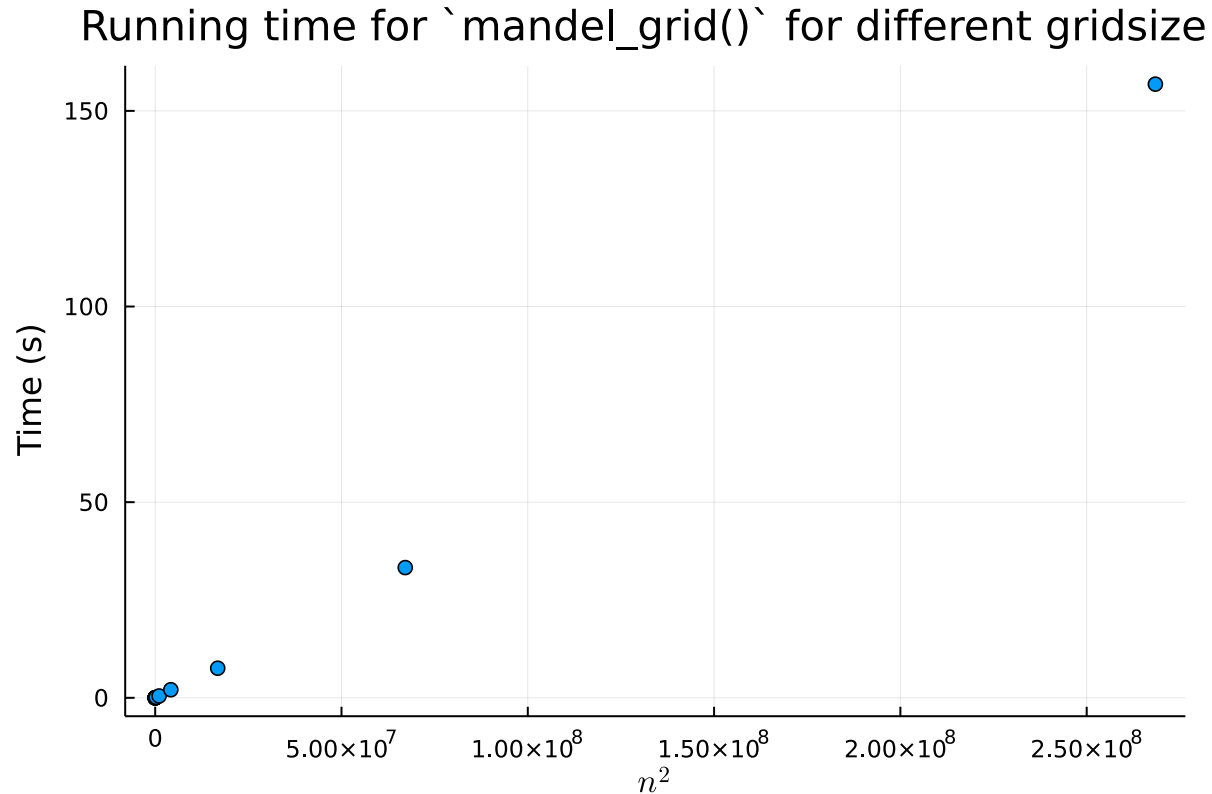
In [40]: using LaTeXStrings

scatter(n_list.^2,

```

```
elapsed,
xlabel=L"n^2",
ylabel="Time (s)",
title="Running time for `mandel_grid()` for different gridsizes",
label=false)
```

Out[40]:



KR7: Discuss the computational complexity of the Mandelbrot function you made based on KR5. What is the best `@time` output to use for this?

As mentioned above the best `@time` macro for this time was `@elapsed`.

Based on the previous two plots, the computational complexity of the Mandelbrot function seems to be linear as a function of actual number of elements in the grid; As a function of the parameter n (length of one side of a square grid), it seems to be at least a polynomial of order 2 based on the magnitudes of n and the coefficients.

In [35]: `using Polynomials`

In [36]: `fit(n_list.^2, elapsed, 1)`

Out[36]: `-0.5419134718089813 + 5.810852347419767e-7 · x`

In [37]: `fit(n_list, elapsed, 2)`

Out[37]: `0.31679617622177886 - 0.0011314560859559033 · x + 6.510779124230988e-7 · x2`

In [38]: `fit(n_list, elapsed, 3)`

Out[38]: `0.0010578525609596418 + 7.233792448958852e-5 · x + 3.935586264020012e-7 · x2 + 1.1369421312555185e-11 · x3`