

Session 4: Fast Function Calls

OBJECTIVE: Compare benchmark times of different implementation of functions that can be expressed as a recursion relation.

KR1: Benchmarked at least two(2) different implementation of the same function or process (e.g. raising each element of an array to some power p , random array may be used) that utilizes some parameter that can be considered a constant or declared globally. Typical methods: (1) Global variable, (2) Constant global variable, and (3) Named parameter variable.

For this session, I will primarily follow the examples in the textbook

```
In [1]: using BenchmarkTools
```

```
In [2]: p = 2

function pow_array(x::Vector{Float64})
    s = 0.0
    for y in x
        s = s + y ^ p
    end
    return s
end
```

```
Out[2]: pow_array (generic function with 1 method)
```

```
In [3]: t = rand(100000)
        @btime pow_array($t)
```

```
Out[3]: 5.851 ms (300000 allocations: 4.58 MiB)
        33260.95810623382
```

```
In [4]: @code_warntype pow_array(t)
```

Variables

```
#self#::Core.Const(pow_array)
x::Vector{Float64}
@_3::Union{Nothing, Tuple{Float64, Int64}}
s::Any
y::Float64
```

Body::Any

```
1 -      (s = 0.0)
|      %2 = x::Vector{Float64}
|          (@_3 = Base.iterate(%2))
|      %4 = (@_3 === nothing)::Bool
|      %5 = Base.not_int(%4)::Bool
|      goto #4 if not %5
2 ... %7 = @_3::Tuple{Float64, Int64}::Tuple{Float64, Int64}
|      (y = Core.getfield(%7, 1))
|      %9 = Core.getfield(%7, 2)::Int64
|      %10 = s::Any
|      %11 = (y ^ Main.p)::Any
|          (s = %10 + %11)
|          (@_3 = Base.iterate(%2, %9))
|      %14 = (@_3 === nothing)::Bool
```

```

└─ %15 = Base.not_int(%14)::Bool
    goto #4 if not %15
3 ─ goto #2
4 ... return s

```

In [5]:

```

const p2 = 2
function pow_array2(x::Vector{Float64})
    s = 0.0
    for y in x
        s = s + y^p2
    end
    return s
end

```

Out[5]: pow_array2 (generic function with 1 method)

Speedup using a constant is very significant when setting p as a constant!

In [6]:

```
@btime pow_array2($t)
```

```

111.183 μs (0 allocations: 0 bytes)
33260.95810623382

```

Out[6]:

In [7]:

```
@code_warntype pow_array2(t)
```

Variables

```

#self#::Core.Const{pow_array2}
x::Vector{Float64}
@_3::Union{Nothing, Tuple{Float64, Int64}}
s::Float64
y::Float64

```

Body::Float64

```

1 ─ (s = 0.0)
└─ %2 = x::Vector{Float64}
    (@_3 = Base.iterate(%2))
    %4 = (@_3 === nothing)::Bool
    %5 = Base.not_int(%4)::Bool
    goto #4 if not %5
2 ... %7 = @_3::Tuple{Float64, Int64}::Tuple{Float64, Int64}
    (y = Core.getfield(%7, 1))
    %9 = Core.getfield(%7, 2)::Int64
    %10 = s::Float64
    %11 = (y ^ Main.p2)::Float64
    (s = %10 + %11)
    (@_3 = Base.iterate(%2, %9))
    %14 = (@_3 === nothing)::Bool
    %15 = Base.not_int(%14)::Bool
    goto #4 if not %15
3 ─ goto #2
4 ... return s

```

In [8]:

```

function pow_array3(x::Vector{Float64})
    return pow_array_inner(x, p)
end

function pow_array_inner(x, pow)
    s = 0.0
    for y in x
        s = s + y ^ pow
    end

```

```
    return s
end
```

Out[8]: pow_array_inner (generic function with 1 method)

In [9]: @btime pow_array3(\$t)

111.357 μ s (1 allocation: 16 bytes)
Out[9]: 33260.95810623382

KR2: Replicated the naive implementation of the polynomial in the textbook.

In [10]:

```
function poly_naive(x, a...)
    p=zero(x)
    for i = 1:length(a)
        p = p + a[i] * x^(i-1)
    end
    return p
end
```

Out[10]: poly_naive (generic function with 1 method)

In [11]: f_naive(x) = poly_naive(x, 1, 2, 3, 4, 5, 6, 7, 8, 9)

Out[11]: f_naive (generic function with 1 method)

In [12]:

```
x = 3.5
bench_naive = @benchmark f_naive($x)
```

Out[12]: BenchmarkTools.Trial: 10000 samples with 709 evaluations.

Range (min ... max):	177.898 ns ... 2.209 μ s	GC (min ... max):	0.00% ... 90.85%
Time (median):	193.292 ns	GC (median):	0.00%
Time (mean \pm σ):	211.444 ns \pm 67.260 ns	GC (mean \pm σ):	0.44% \pm 2.01%



Memory estimate: 32 bytes, allocs estimate: 2.

KR3: Replicated the naive implementation of the Horner's method for the same polynomial.

In [13]:

```
function poly_horner(x, a...)
    b=zero(x)
    for i = length(a):-1:1
        b = a[i] + b * x
    end
    return b
end
```

Out[13]: poly_horner (generic function with 1 method)

In [14]:

```
f_horner(x) = poly_horner(x, 1,2,3,4,5,6,7,8,9)
bench_horner = @benchmark f_horner($x)
```

Out[14]: BenchmarkTools.Trial: 10000 samples with 1000 evaluations.

```

Range (min ... max): 4.884 ns ... 4.262 μs | GC (min ... max): 0.00% ... 0.00%
Time (median): 5.293 ns | GC (median): 0.00%
Time (mean ± σ): 8.030 ns ± 54.682 ns | GC (mean ± σ): 0.00% ± 0.00%

```



Memory estimate: 0 bytes, allocs estimate: 0.

KR4: Replicated the macro implementation of the Horner's method of the same polynomial.

```

In [15]: macro horner(x, p...)
          ex = esc(p[end])
          for i = length(p)-1:-1:1
              ex = :(muladd(t, $ex, $(esc(p[i]))))
          end
          Expr(:block, :(t = $(esc(x))), ex)
        end

```

Out[15]: @horner (macro with 1 method)

```

In [16]: f_horner_macro(x) = @horner(x, 1,2,3,4,5,6,7,8,9)

```

Out[16]: f_horner_macro (generic function with 1 method)

```

In [17]: bench_macro = @benchmark f_horner_macro($x)

```

```

Out[17]: BenchmarkTools.Trial: 10000 samples with 1000 evaluations.
Range (min ... max): 0.046 ns ... 26.623 ns | GC (min ... max): 0.00% ... 0.00%
Time (median): 0.049 ns | GC (median): 0.00%
Time (mean ± σ): 0.056 ns ± 0.282 ns | GC (mean ± σ): 0.00% ± 0.00%

```



Memory estimate: 0 bytes, allocs estimate: 0.

KR5: Table showing how many minutes will the function evaluations in both KR3 and KR4 be reduced if KR2 requires 24hours of runtime.

```

In [18]: method_names = ["Naive", "Horner", "Macro"];
          speedup = [median(i.times) for i in [bench_naive, bench_horner, bench_macro]] / median(ben
          minutes = speedup * 1440;;

```

```

In [19]: using DataFrames
          table = DataFrame("Method"=>method_names,"Speedup" => speedup, "Runtime" => minutes);

          print(table)

```

3x3 DataFrame

Row	Method	Speedup	Runtime
	String	Float64	Float64
1	Naive	1.0	1440.0
2	Horner	0.0273834	39.4322
3	Macro	0.000253503	0.365044

The speedup using the Horner method is very impressive, from a 1 day runtime to just 40 minutes and when using a macro it actually just takes a few seconds. I actually didn't expect that much of a speedup especially for the macro and I'll always consider this if ever I run into future bottlenecks.

In []: